

XJCO3221 Parallel Computation

Peter Jimack and David Head

University of Leeds

Lecture 2: Introduction to shared memory parallelism (SMP)

Previous lecture

In the introductory lecture we saw:

- Why **technological limitations** have led to multi-core CPUs.
- Parallel architectures also present in high-performance **clusters**, and **graphics processing units** (GPUs).
- Some general concepts:
 - **Concurrency** (more general than parallelism).
 - **Shared *versus* distributed memory**.
 - Potential performance issues related to **communication**.
 - **Flynn's taxonomy**.

This lecture

This lecture is the first of six on **shared memory parallelism**, relevant to **multi-core CPUs**.

- The hardware **architecture**, including the memory cache.
- Processes *versus* threads and the thread **scheduler**.
- Languages and frameworks suitable for these systems.
- How to set up and run OpenMP.

Multi-core CPUs

A **core** is a single **processing unit** that executes instructions:

- Has components that *fetch, decode etc.* instructions.
- **Functional units** for integer and floating point operations.
- Other features such as **instruction level parallelism**.

As its name suggests, **multi-core** processors contain more than one such unit.

- **MIMD** in Flynn's taxonomy (*single cores are SISD*).
- Most common now are **dual core**, **quad core** and **octa core**.
- High-performance chips can have many more, e.g. SW26010 (used in China's Sunway TaihuLight supercomputer) has 260.

Simultaneous multithreading

Some chips employ **simultaneous multithreading**¹:

- Two (or more) threads run on the same core.
- If one thread stops execution (e.g. to wait for memory access), the other takes over.

Appears as two **logical processors** to the programmer, and only requires 5% increase in chip area.

- Performance improvements only 15%-30%.²

When interrogating a framework for the maximum number of available threads, you may get **more** than the number of cores.

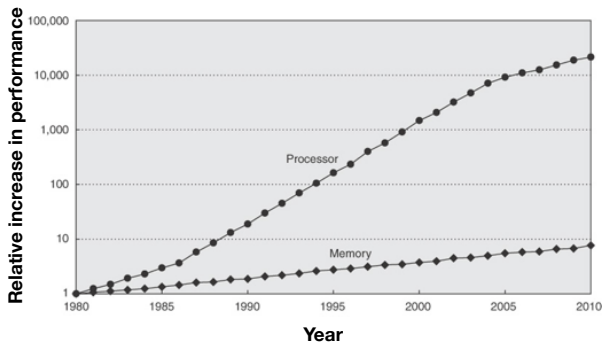
¹Known as **hyperthreading** on Intel chips.

²Rauber and Rünger, *Parallel Programming* 2nd ed. (Springer, 2013).

The processor-memory gap

Memory access rates are increasing far slower than processor performance (taking into account number of cores):

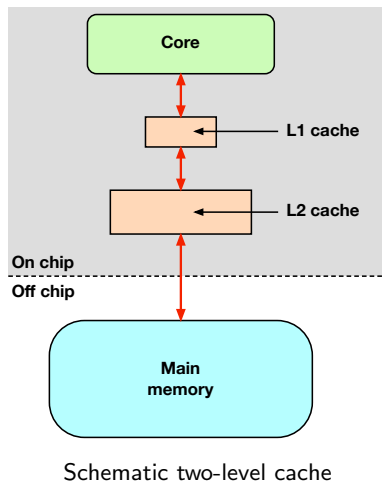
- This is the **processor-memory gap**.



Hennessy and Patterson, *Computer Architecture: A Quantitative Approach* (Morgan Kauffman, 2006).

Single-core memory caches: A reminder

- **Small, fast, on-chip** memory.
- Accessing main memory returns a **line** to the cache (e.g. 64 bytes).
- Subsequent accesses return the cache data (a **cache hit** - *fast*), or from main memory (a **cache miss** - *slow*).
- Multiple caches **levels** (e.g. L1, L2, L3) arranged **hierarchically**.



Multi-core memory caches

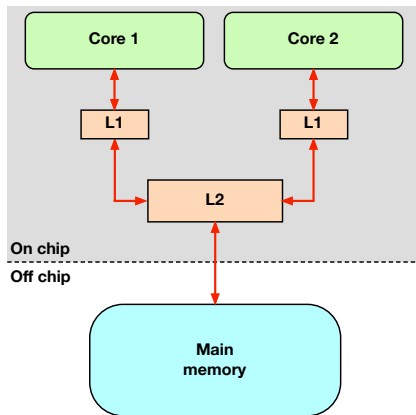
Different manufacturers choose different ways to incorporate caches into multi-core designs.

Often use **hierarchy**:

- Each core has its own L1 cache.
- Share higher level caches.

For e.g. quad cores:

- L1 for each core.
- L2 for pairs.
- L3 for all cores.



Cache coherency

- 1 Core 1 reads an address x , resulting in a line in its L1.
- 2 Core 2 does the same, resulting in a line in its L1.
- 3 Core 1 changes the value of x in its L1.
- 4 Core 2 reads x from its L1, which **still has the old value**.

Maintaining consistent memory views for all cores is known as
cache coherency

A common way to maintain **cache coherency** is **snooping**:

- The *cache controller* **detects writes** to caches, and updates higher-level caches.

False sharing

Maintaining cache coherency incurs a performance loss.

- If two cores repeatedly write to the same memory location, the higher level caches will be constantly updated.

However, if the cores write to nearby **but different** memory locations **on the same cache line**, updates will still occur.

- *i.e.* hardware **performance loss** with no need.

This unnecessary cache coherency is known as **false sharing**

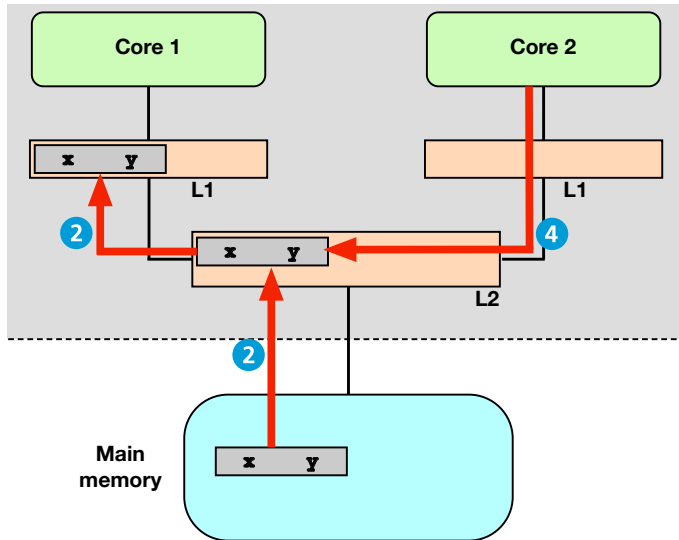
Potential benefit of cache sharing

It is also possible for multiple cores to benefit from shared caches:

- ① Core 1 reads an address x from main memory.
- ② A line including x is read into L2, and the L1 for Core 1.
- ③ Core 2 now tries to access an address y that is 'near to' x .
- ④ If y is on the line just copied into L2, Core 2 will **not need to access main memory**.

It is therefore possible for fewer accesses to main memory overall, compared to the equivalent serial code.

This can result in parallel speed up **more than the number of cores** (known as superlinear speedup; cf. Lecture 4).



Processes *versus* threads

Control flows can be either **processes** or **threads**:

Processes:

- **Executable program** plus all required information.
- Register, stack and heap memory, with own address space.
- Explicit communication between processes (*via* sockets).
- **Expensive** to generate (large heap memory).

Threads:

- Threads of one process **share** its address space.
- Implicit communication *via* this **shared memory**.
- **Cheap** to generate (no heap memory).

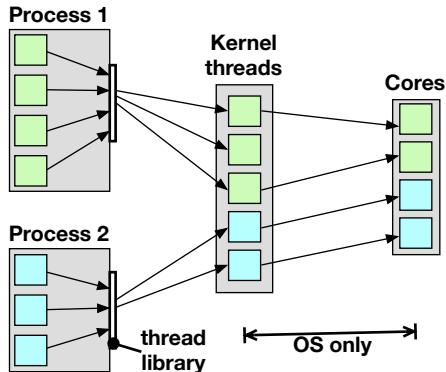
Kernel *versus* user-level threads

The threads that execute on the core(s) are **kernel threads**.

- Only the OS has direct control over kernel threads.

Programmers instead generate **user-level threads**.

- Managed by a **thread library**.
- Mapped to kernel threads by the OS **scheduler**.



Thread programming

The choice of programming framework/API/library/etc. **must** be suitable for the architecture on which it will run.

- Multi-core CPUs use **SMP** = Shared Memory Parallelism.

It is possible to program user threads directly:

- Java supported threads early on, through the Thread class and Runnable interface.
- The C library pthread implements POSIX threads.
- C++11 has language-level concurrency support.
- Python has a threading library, although need to work around its **global interpreter lock** to exploit multi-cores.

Higher-level threading support

Higher-level options that do not require explicit thread control also exist to reduce development times.

- Java's Concurrency library (in `java.util`).
- The OpenMP standard (this module and next slides).

For C/C++, as well as OpenMP there are [see McCool *et al.* in *Structured Parallel Programming* (Morgan-Kaufman, 2012)].

- **Cilk Plus**.
- **TBB** (Threading Building Blocks).
- **ArBB** (Array Building Blocks).
- **OpenCL**, although primarily used for GPUs.

The first three are not (yet?) widely implemented in compilers.

OpenMP

For the SMP component of this module we will use **OpenMP**.

- Portable standard devised in 1997 and widely implemented in C, C++ and FORTRAN compilers.
- Maintained by the OpenMP Architecture Review Board.
- Currently up to v5.2, although compilers may only support earlier versions

More information available from <http://www.openmp.org>

Compiling C with OpenMP

For this module we will use gcc (GNU Compiler Collection)¹:

- Compile with `-fopenmp`
- Must include `omp.h`

All parallel execution will be undertaken via the Cloud

- We have created a mini-HPC environment using Microsoft Azure.
- You will each get your own INDIVIDUAL account.
- Full instructions for logging into your account will be provided.
- A Linux OS is provided along with the all libraries required for this module:
 - You can run jobs interactively on 2 cores while debugging (can still have more threads!)
 - You can run batch jobs on up to 16 cores via Slurm

¹Easy to install on Macs with homebrew.

helloWorld.c

Code on Minerva: helloWorld.c

```
1 #include <stdio.h>
2 #include <omp.h> // Run-time OpenMP library routines.
3
4 int main()
5 {
6     // The scope after this pragma is in parallel.
7     #pragma omp parallel
8     {
9         // Get this thread number, and the maximum.
10        int threadNum = omp_get_thread_num ();
11        int maxThreads = omp_get_max_threads();
12
13        // Simple message to stdout.
14        printf( "Hello from thread %i of %i!\n", threadNum,
15              maxThreads );
16    }
17    return 0;
18 }
```

Compiling C: Reminder

If the source code is called `helloWorld.c`:

```
gcc -fopenmp -Wall -o helloWorld helloWorld.c
```

Options:

- `-fopenmp` tells compiler to expect OpenMP pragmas.
- `-Wall` turns on all warnings; recommended but not required.
- `-o helloWorld` is the executable name (a.out by default).
- `helloWorld.c` is the source code.
- Sometimes need e.g. `-lm` for the maths library.

#pragma omp parallel

#pragma directives provide information beyond the language.

- All OpenMP pragmas start: `#pragma omp ...`

Here, `#pragma omp parallel` is telling the compiler to perform the next scope (*i.e.* the section of code between the curly brackets, from `{` to `}`) **in parallel**.

- The code inside this scope is run by **multiple threads**.
- Outside of this scope there is only **one thread**.

This is why the `printf` statement is repeated multiple times, **even though it only appears once in code**.

We will look at this in more detail next time.

#include <omp.h>

Include `omp.h` to use OpenMP runtime library routines:

`int omp_max_thread_num():`

- Returns the maximum number of threads.
- Defaults to **hardware concurrency**, e.g. the number of cores.
- May exceed apparent core number with **simultaneous multithreading** (see earlier).

`int omp_get_thread_num():`

- Returns the thread number **within the current scope**.
- $0 \leq \text{omp_get_thread_num}() < \text{omp_max_thread_num}()$

Setting the number of threads

If you don't want to use the default number of threads:

```
void omp_set_num_threads(int):
```

- Will change the number of threads **dynamically**.
- **Can** exceed hardware concurrency.

Alternatively, use shell **environment variables**:

- For bash: `export OMP_NUM_THREADS=<num>`
- Avoids the need to recompile.
- List all environment variables using `env`.
- To see all OMP variables: `env | grep OMP`

Summary and next lecture

Today we have started looking at shared memory parallelism (SMP):

- Relevant to **multi-core CPUs**.
- OS **scheduler** maps **threads** to cores.
- Various languages, frameworks *etc.* support SMP.
- OpenMP is commonly supported by C/C++ compilers.

Next time we will look in more detail at what is actually going on at the thread level, for a more interesting example.