

# XJC03221 Parallel Computation

Peter Jimack

University of Leeds

Lecture 3: Data parallel problems

# Previous lectures

In the last lecture we started looking at **shared memory parallelism** (SMP):

- Relevant to **multi-core CPUs**.
- Separate processing units (cores) share some levels of **memory cache**.
- Various frameworks for programming SMP systems.
- Widely-implemented standard: **OpenMP**.

# Today's lecture

Today we are going to look at a some actual problems.

- Examples of a **data parallel** problems, where the same operation is applied to multiple data elements.
- Also known as a **map**<sup>1</sup>.
- **Multi-threading** solution employs a **fork-join** pattern.
- How to parallelise **nested loops**.
- Parallel code can be **non-deterministic**, even when the serial code is **deterministic**.

---

<sup>1</sup>McCool *et al.*, *Structured parallel programming* (Morgan-Kaufman, 2012).

# Vector addition

An **n-vector** **a** can be thought of as an **array** of  $n$  numbers:  
 $\mathbf{a} = (a_1, a_2, \dots, a_n)$ .

If two vectors **a** and **b** are the same size, they can be added to generate a new  $n$ -vector **c**:

$$\begin{array}{cccccc}
 \mathbf{a} = ( & a_1, & a_2, & a_3, & \dots, & a_n & ) \\
 + & + & + & + & & + & \\
 \mathbf{b} = ( & b_1, & b_2, & b_3, & \dots & b_n & ) \\
 \downarrow & \downarrow & \downarrow & \downarrow & & \downarrow & \\
 \mathbf{c} = ( & c_1, & c_2, & c_3 & \dots, & c_n & )
 \end{array}$$

Or:

$$c_i = a_i + b_i \quad , \quad i = 1 \dots n.$$

# Serial vector addition

Code on Minerva: `vectorAddition_serial.c`

```
1  #define n 100
2
3  int main()
4  {
5      float a[n], b[n], c[n];
6
7      ... // Initialise a[n] and b[n]
8
9      int i;
10     for( i=0; i<n; i++ )
11         c[i] = a[i] + b[i];
12
13     return 0;
14 }
```

Note that indices usually start at 0 for most languages, but 1 for the usual mathematical notation (also FORTRAN, MATLAB).

# Vector addition in parallel

Code on Minerva: `vectorAddition_parallel.c`

Add `#pragma omp parallel` for just before the loop:

```
1  #define n 100
2
3  int main()
4  {
5      float a[n], b[n], c[n];
6
7      ... // Initialise a[n] and b[n]
8
9      int i;
10     #pragma omp parallel for
11     for( i=0; i<n; i++ )
12         c[i] = a[i] + b[i];
13
14     return 0;
15 }
```


**This only parallelises this one loop, not any later ones!**

# Fork-and-join

When the executable reaches `#pragma omp parallel for`, it spawns multiple **threads**.

- Each thread computes **part** of the loop.
- The extra threads are **destroyed** at the end of the loop.

This is known as a **fork-join** construct:



## Example: Four threads in total

Pseudocode for the **main** thread:

```
1 // Main thread starts in serial
2 // Initialise arrays a, b; allocate c.
3 ...
4 // REACHES #pragma omp parallel for
5 // FORK: Create three new threads.
6 worker1 = fork(...);
7 worker2 = fork(...);
8 worker3 = fork(...);
9
10 // Perform 1/4 of the total loop.
11 for( i=0; i<n/4; i++ )
12     c[i] = a[i] + b[i];
13
14 // JOIN: Wait for other threads to finish.
15 worker1.join();
16 worker2.join();
17 worker3.join();
18
19 // Continue in serial after the loop
```



## Worker thread 1:

```
1 // CREATED BY MAIN ('fork')
2 // Perform second 1/4 of loop.
3 for( i=n/4; i<n/2; i++ ) c[i] = a[i] + b[i];
4 // FINISH ('join')
```

## Worker thread 2:

```
1 // CREATED BY MAIN ('fork')
2 // Perform third 1/4 of loop.
3 for( i=n/2; i<3*n/4; i++ ) c[i] = a[i] + b[i];
4 // FINISH ('join')
```

## Worker thread 3:

```
1 // CREATED BY MAIN ('fork')
2 // Perform final 1/4 of loop.
3 for( i=3*n/4; i<n; i++ ) c[i] = a[i] + b[i];
4 // FINISH ('join')
```

# Notes

The four threads are **not** being executed one after the other:

- Each thread runs **concurrently**, hopefully on separate cores, *i.e.* in **parallel**.
- Cannot be understood in terms of serial programming concepts.

Each thread performs the **same** operations on **different** data.

- Would be **SIMD** in Flynn's taxonomy, except this is implemented **in software** on a MIMD device.

Have assumed  $n$  is divisible by the number of threads for clarity.

- Generalising to arbitrary  $n$  is not difficult, but obscures the parallel aspects.

## #pragma omp parallel for

The total loop range was evenly divided between all threads.

- Happens as soon as `#pragma omp parallel for` reached.
- The **trip count** (*i.e.* loop range) **must** be known at the **start** of the loop.
- The start, end and stride must be **constant**.
- Cannot break from the loop.
- **Cannot apply to 'while...do' or 'do...while' loops**

# Data parallel and embarrassingly parallel

This is an example of a **data parallel problem** or a **map**:

- Array elements distributed evenly over the threads.
- Same operation performed on all elements.
- Suitable for the SIMD model.

In fact, this example is so straightforward to parallelise that is also sometimes referred to as an **embarrassingly parallel problem**.


- Easy to get working **correctly** in parallel.
- May still be a challenge to achieve good parallel **performance**.

# Mandelbrot set generator

Code on Minerva: `Mandelbrot.c`, `makefile`

Classic computationally intensive problem in two dimensions that used to be used as a **benchmark** for processor speeds:

- Loops over **pixels**, *i.e.* a **two dimensional, nested** double loop.
- Colour of each pixel calculated **independently** of all other pixels.
- Each colour calculation requires many **floating point operations**.



## Code snippet

The part of the code that interests us here is shown below:

```
1 // Change the colour arrays for the whole image.
2 int i, j;
3 for( j=0; j<numPixels_y; j++ )
4     for( i=0; i<numPixels_x; i++ )
5     {
6         // Set the colour of pixel (i,j), i.e. modify the values
7         // of red[i][j], green[i][j], and/or blue[i][j].
8         setPixelColour( i, j );
9     }
```

Note the i-loop is nested inside the j-loop.

The graphical output is performed in OpenGL/GLFW. Since including and linking is different between Linux and Macs, a simple makefile has been provided.

# What setPixelColour does

Purely for background interest, here's how the colours are calculated:

- 1 Each **pixel**  $i, j$  is converted to **floating point numbers**  $c_x, c_y$ , both in the range -2 to 2.
- 2 Two other floats  $z_x$  and  $z_y$  are initialised to zero.
- 3 The following iteration<sup>1</sup> is performed until  $z_x^2 + z_y^2 \geq 4$ , or a maximum number of iterations `maxIters` is reached:

$$(z_x, z_y) \rightarrow (z_x^2 - z_y^2 + c_x, 2z_xz_y + c_y)$$

- 4 The colour is selected based on the number of iterations.

---

<sup>1</sup>More concisely represented as **complex numbers**  $c$  and  $z$  [with e.g.  $z_x = \Re(z)$ ], then the iteration is just  $z \rightarrow z^2 + c$ .

## Parallel Mandelbrot: First attempt

Parallelise **only** the **inner** loop.

```
1 int i, j;  
2 for( j=0; j<numPixels_y; j++ )  
3     #pragma omp parallel for  
4     for( i=0; i<numPixels_x; i++ )  
5     {  
6         setPixelColour( i, j );  
7     }
```

This works, but may be slower than serial (*check on your system*).

Multiple possibilities for this:

- The **fork-join** is **inside** the j-loop, so threads are created and destroyed `numPixels_y` times, which incurs an **overhead**.
- This problem suffers from poor **load balancing**; see later.



## Parallel Mandelbrot: Second attempt

Parallelise only the **outer** loop, so there is only a single **fork** event and a single **join** event.


```
1 int i, j;  
2 #pragma omp parallel for  
3 for( j=0; j<numPixels_y; j++ )  
4     for( i=0; i<numPixels_x; i++ )  
5     {  
6         setPixelColour( i, j );  
7     }
```

This is faster ... but **wrong**!

- A **distorted** image results.
- The distortion is **different each time** the program is executed.

The **same** variable `i` for the inner loop counter is being **updated by all threads**:

- When one thread completes a calculation, it increments `i`.
- Therefore other threads will skip at least one pixel.
- Threads do **not** calculate the full line of pixels.



## Parallel Mandelbrot: Third attempt

Make the inner loop variable **i** **private** to each thread:

```
1 int j;  
2 #pragma omp parallel for  
3 for( j=0; j<numPixels_y; j++ )  
4 {  
5     int i;  
6     for( i=0; i<numPixels_x; i++ )  
7     {  
8         setPixelColour( i, j );  
9     }  
10 }
```

...or (for compilers following the C99 standard):

```
1 #pragma omp parallel for  
2 for( int j=0; j<numPixels_y; j++ )  
3     for( int i=0; i<numPixels_x; i++ )  
4     {  
5         setPixelColour( i, j );  
6     }
```

# The private clause

A third way to solve this is to use OpenMP's private **clause**:

```
1 int i, j;  
2 #pragma omp parallel for private(i)  
3 for( j=0; j<numPixels_y; j++ )  
4     for( i=0; i<numPixels_x; i++ )  
5     {  
6         setPixelColour( i, j );  
7     }
```

- Creates a **copy** of *i* for each thread.
- Multiple variables may be listed, e.g. `private(i,a,b,c)`

The code now works ... but is no faster than serial!

- The primary overhead is poor **load balancing**. We will look at this next lecture briefly, and detail in Lecture 13.

# The collapse clause

The collapse clause replaces 2 or more nested loops with a single loop, at the expense of additional internal calculations.

```
1 #pragma omp parallel for collapse(2)
2 for( int j=0; j<numPixels_y; j++ )
3     for( int i=0; i<numPixels_x; i++ )
4         setPixelColour( i, j );
```

is equivalent to (but more readable than)


```
1 #pragma omp parallel for
2 for( int k = 0; k < numPixels_x * numPixels_y; k++ )
3     {
4         int
5             i = k % numPixels_x,
6             j = k / numPixels_x;
7         setPixelColour( i, j );
8     }
```

This is principally intended for **short** loops that cannot be equally distributed across all threads.


# Determinism and non-determinism

Notice that the incorrect images were slightly different each time:

e.g. 1.



e.g. 2.



The pixels plotted depend on the order in which threads update the shared variable `i`, which depends on the thread **scheduler**.

- Will be influenced by factors outside our control.
- e.g. the various **background tasks** that every OS must run.

Our serial code was **deterministic**, *i.e.* produced the same results each time it was run.

By contrast, our (incorrect) parallel code was **non-deterministic**.

Often this is the result of an error, but can sometimes be useful:

- Some algorithms, often in science and engineering, do not care about non-deterministic errors **as long as they are small**.
- Strictly imposing determinism may result in additional overheads and performance loss.

However, for this module we will try to develop parallel algorithms whose results match that of the serial equivalent.

## Summary and next lecture

Today we have look at **data parallel** problems or **maps**, where the same operation is applied to multiple data members.

- Distribute data evenly across threads.
- Sometimes referred to as **embarrassingly parallel**.

In two lectures time we will start looking at more complex problems for which the calculations on different threads are **not** independent.

Before then, we need to learn the vocabulary of parallel theory, which is the topic of next lecture.