

XJCO3221 Parallel Computation

Peter Jimack

University of Leeds

Lecture 5: Loop parallelism and data races

Previous lectures

In Lecture 3 we saw how two problems could be parallelised:

- 1 **Vector addition**, where two one-dimensional arrays were added together element-by-element.
- 2 **Mandelbrot set**, where pixel colours were calculated to generate a 2D image.

However, neither of these problems have any **data dependencies**.

- 1 Each vector element was calculated **independently** of other elements.
- 2 Each pixel colour was calculated **independently** of the others.

This lecture

In this lecture we will look how to parallelise problems that **do** have data dependencies.

- Can lead to **data races** on shared memory systems.
- Behaviour then becomes **non-deterministic**.
- May require algorithm re-design to remove dependencies.

We will then look at three examples of **loop parallel** problems and how their dependencies can be resolved.

Example of a data race

Consider the following pseudocode¹ for two concurrent threads, where each thread accesses the same variable x . $x=0$ at the start of the code segment.

Thread 0:

```
a = x;  
a += 1.0;  
x = a;
```

Thread 1:

```
b = x;  
b += 2.0;  
x = b;
```

What value does x take at the end?

¹From §2.6 of McCool *et al.*, *Structured parallel programming* (Morgan-Kaufman, 2012).

Non-determinism

The result may differ each time the program is executed.

- An example of **non-determinism** [*Lecture 3*].

The scheduler runs threads depending on various factors, including other applications and OS tasks.

- Cannot predict which thread is **launched** first.
- The OS may **suspend** threads to make way for other tasks (*pre-emptive multitasking*).
- The instructions may become **interleaved**.

Interleaved instructions (example)

Recall $x=0$ initially.

```
1  a  = x;    // Thread 0: a now 0
2  a += 1;    // Thread 0: a now 1
3  b  = x;    // Thread 1: b now 0
4  x  = a;    // Thread 0: x now 1
5  b += 2;    // Thread 1: b now 2
6  x  = b;    // Thread 1: x now 2
```

In this example, $x=2$ at the end.

- Possible to get $x=1$ or $x=3$ by different interleaving of instructions (*check left as exercise*).

Race conditions

This is known as a **data race** or a **race condition**.

- Result of calculation depends on which thread reaches its instructions first.
- **Non-deterministic**, which is usually undesirable.

Only an issue for **shared memory**.

- If each thread had its own x , there would be no race.
- In this example, if each thread had its own x , $x=1$ for thread 0 and $x=2$ for thread 1 at the end, regardless of any interleaving.

Read-only does not lead to a data race

For a race condition to arise, **at least one** thread must **write** to `x`.

- No race if all threads just **read** `x`.

There is no race for the following example, as both threads only read `x`:

Thread 0:

```
a = x;  
a += 1.0;
```

Thread 1:

```
b = x;  
b += 2.0;
```

For this example, `x=0` (and `a=1` and `b=2`) at the end.

Sequential consistency

Have assumed each thread executes its instructions in order.

- *i.e.* have assumed **sequential consistency**.

Compilers often **rearrange instructions** to improve performance.

- *e.g.* bring forward memory accesses, combine operations *etc.*
- The result is the same **in serial**.
- However, **multithreading** can confuse compilers.
- In the original example, this means we can even get $a=b=0!$

The volatile keyword

You may read that the way to solve this is to declare variables as `volatile` (in C/C++). However, this is only partially correct¹.

- `volatile` is for **special memory**, such as that read/written by external device (*memory-mapped I/O*).
- The way compilers handle such variables is **not guaranteed** to work for multithreading.

If this might be an issue, should use features **specific** to concurrent programming.

- e.g. memory fences, `std::atomic<>` in C++11 etc.

Will come to atomics next lecture and Lecture 18.

¹S. Meyers, *Effective modern C++* (O'Reilly, 2015).

Loop parallelism

Often we are required to parallelise **loops**.

- Known as **loop parallelism**.
- If there are **data dependencies**, may have implicit **data races** within the loop.

There is no **systematic** way to parallelise loops.

- How to remove a dependency depends on context.
- Consider extra resources or loops to reach the correct solution.

For the remainder of this lecture, will give examples of loops with data dependencies, and how to overcome them.

Example 1: Redundant variable(s)

Consider the following serial code:

```
1 float temp, a[n], b[n], c[n];  
2 ... // Initialise arrays b and c  
3  
4 int i;  
5 for( i=0; i<n; i++ )  
6 {  
7     temp = 0.5f*( b[i] + c[i] );  
8     a[i] = temp;  
9 }
```

Here, temp is being used as a temporary variable.

- Sometimes useful to make (more complex) code easier to read.

Need to make temp a private (or local) variable:

```
1 #pragma omp parallel for
2 for( i=0; i<n; i++ )
3 {
4     float temp = 0.5f*( b[i] + c[i] );
5     a[i] = temp;
6 }
```

Can also use OpenMP's private clause:

```
1 #pragma omp parallel for private(temp)
2 for( i=0; i<n; i++ )
3 {
4     temp = 0.5f*( b[i] + c[i] );
5     a[i] = temp;
6 }
```

cf. the inner loop counter in Lecture 3's Mandelbrot set example.

Example 2: Shift dependency

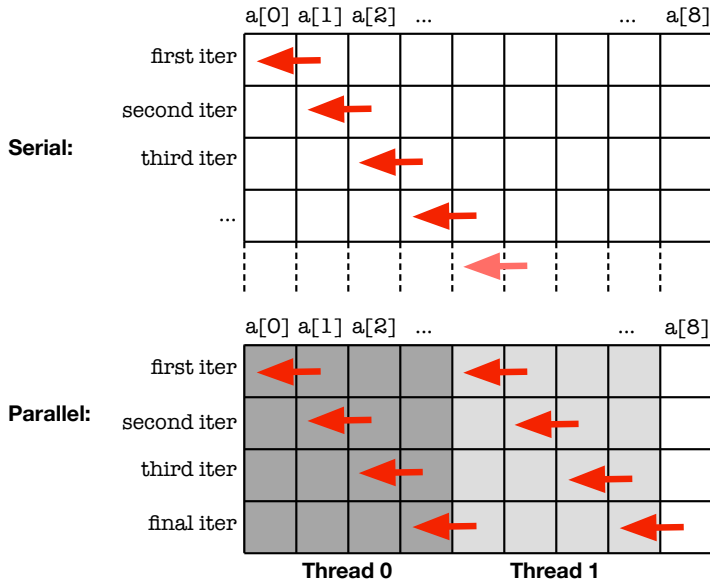
Code on Minerva: `shiftDependency.c`

Consider a shift dependency:

```
1 float a[n];  
2 ... // Initialise array a  
3  
4 int i;  
5 for( i=0; i<n-1; i++ )  
6     a[i] = a[i+1];
```

Naive parallelisation does not *quite* work:

```
1 #pragma omp parallel for  
2 for( i=0; i<n-1; i++ )  
3     a[i] = a[i+1];
```



A solution here is to **copy** the array **a** **before** the loop:

```
1 float atemp[n];  
2  
3 #pragma omp parallel for  
4 for( i=1; i<n; i++ )  
5     atemp[i] = a[i];  
6  
7 #pragma omp parallel for  
8 for( i=0; i<n-1; i++ )  
9     a[i] = atemp[i+1];
```

This comes at the expense of **additional resources**:

- **Memory** for the array atemp.
- **CPU time** to copy a to atemp.

Examples of **parallel overheads**.

Example 3: Red-black Gauss-Seidel

Code on Minerva: `redBlackGaussSeidel.c`

Finally consider this:

```
1 for( i=1; i<n-1; i++ )  
2   a[i] = 0.5f * ( a[i+1] + a[i-1] );
```

Each element takes the average of the elements either side of it.

- Can be used to **smooth** vector **a**, e.g. in image transformations ('blurring') with a 2D nested loop over pixels.
- Also arises in numerical computation - the **diffusion** or **heat equation** solved using the **Gauss-Seidel** method.

We **could** make a copy atemp as before.

- In numerical computation, this is the **Jacobi method**.

However, this is undesirable in some situations:

- Typically repeat the loop until 'reaching' the solution, which is faster (*takes fewer iterations*) when not using a copy.

Instead we consider a **modified serial** variant which is more amenable to parallelisation.

- Known as **red-black Gauss-Seidel**, as it bears a resemblance to red and black squares on a chessboard (in 2D).

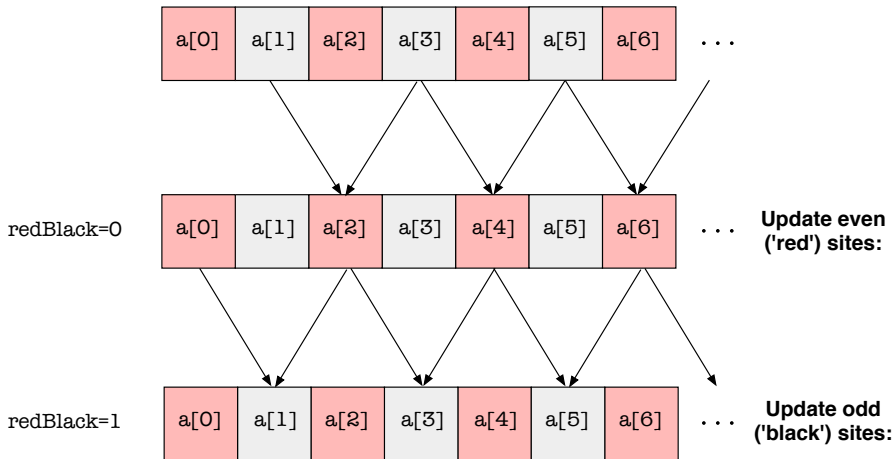
Update **even** elements first, then **odd** elements¹:

```
1 int redBlack;
2 for( redBlack=0; redBlack<2; redBlack++ )
3     for( i=1; i<n-1; i++ )
4         if( i%2 == redBlack )
5             {
6                 a[i] = 0.5f * ( a[i-1] + a[i+1] );
7             }
```

The outer redBlack loop only loops over 2 values, 0 or 1.

- When redBlack=0, only the elements of a[i] with i **even** are updated.
- Similarly, only the **odd** are updated when redBlack=1.

¹Recall i%2 gives the remainder of division by 2, so e.g. i%2==0 if i is even.



Note that for each loop, the calculations are **now independent**.

- We have **removed the dependency**, albeit by slightly changing the serial algorithm.

Now clear how to parallelise:

```
1 for( redBlack=0; redBlack<2; redBlack++ )
2     #pragma omp parallel for
3     for( i=1; i<n-1; i++ )
4         if( i%2 == redBlack )
5             {
6                 a[i] = 0.5f * ( a[i-1] + a[i+1] );
7             }
```

There are no dependencies **within** the i-loop, because $a[i-1]$ and $a[i+1]$ were/will be updated in the other redBlack loop.

The 'best' parallel algorithm?

Notice that we changed Gauss-Seidel to the red-black variant to make it more efficient in parallel.

- For Gauss-Seidel, this is the standard parallel variant.

Perfectly acceptable to do this if:

- 1 The solution reached is still 'correct.'
- 2 Parallel scaling is good (for large arrays).

As a general observation, the 'best' parallel algorithm need **not** be directly related to the 'best' serial algorithm.

Summary and next lecture

This lecture we have seen how multiple threads with at least one writing to shared memory can lead to **data races**.

- Outcome is not predictable (*non-deterministic*).
- Can be a challenge to remove **data dependencies** from loops.
- May require additional resources not present in the serial version, *i.e.* **parallel overheads**.

The next lecture will look at a way to **synchronise** threads within a parallel program and apply it to a linked list.