

# XJC03221 Parallel Computation

Peter Jimack

University of Leeds

Lecture 7: Lock and mutexes

## Previous lecture

In the last lecture we saw how critical regions of code could be **serialised**:

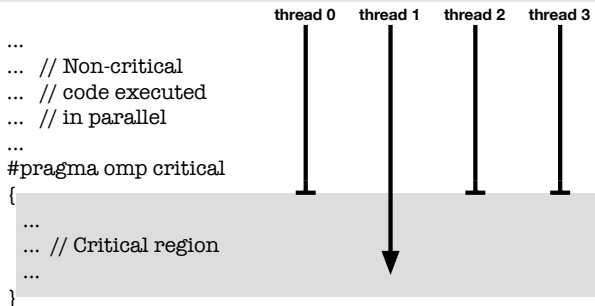
- Only one thread can enter the region at a time, a form of **coordination**.
- Avoids **data races**.
- Can incur a significant **performance penalty**.
- Implemented in OpenMP as `#pragma omp critical`
- Single arithmetic instructions can be optimised by using **atomic** instructions (`#pragma omp atomic`).

# This lecture

For this final lecture on shared memory parallelism, we will look at what is going on 'behind the scenes'.

- Thread coordination performed using **locks**, sometimes referred to as **mutexes**.
- Locks can control access to **data structures**.
- Multiple locks can improve **performance** of memory access.
- However, multiple locks can give rise to **deadlock**.

## Recap: Critical regions



- Instructions before `#pragma omp critical` executed **concurrently** (e.g. if in a parallel loop).
- Instructions in the scope (`{` to `}`) only executed by **one thread at a time**.
- Other threads blocked from entering; they are **idle**.

# Thread coordination with locks

This **synchronisation** is performed using a **lock**:

- Single lock for the critical region.
- Can be in one of two states: **locked** and **unlocked**.
- The first thread to reach the critical region **locks** it.
  - Also known as **acquiring the lock**.
  - This thread is said to be the lock's **owner**.
- No other threads can enter the region ('acquire the lock') until it becomes unlocked.
- The owning thread **unlocks** (or **releases**) it when leaving the region, allowing another thread to take over ownership.

# Critical region using a lock

## OpenMP:

```
1 // Multiple threads executing concurrently,  
2 // (e.g. parallel loop).  
3  
4  
5 #pragma omp critical  
6 {  
7     ...  
8     ... // Critical code.  
9     ...  
10 }
```

## Lock pseudocode:

```
1 // All threads access a  
2 // single lock object.  
3 lock_t regionLock;  
4  
5 regionLock.lock();  
6  
7 ...  
8 ... // Critical code.  
9 ...  
10 regionLock.unlock();
```

`regionLock.lock()` does not return until the thread has **acquired** the lock; it is said to be **blocking**.

# Implementations of locks

Most parallel APIs support **locks**, although they are sometimes called **mutexes** as they control mutual exclusion:

- Java's Lock interface (in `java.util.concurrent.locks`).
- `std::mutex` in C++11 (in `<mutex>`).
- `pthread_mutex_t` in the pthreads library (C/C++).

When implemented as classes, they are typically **opaque**:

- The user does not have access to instance variables or details of the implementation.

# Locks in OpenMP

OpenMP also supports locks.

```
1 #include <omp.h>
2
3 // Initialise lock (opaquely).
4 omp_lock_t regionLock;
5 omp_init_lock(&regionLock);
6 ...                               // (in parallel).
7 omp_set_lock(&regionLock);        // LOCK.
8 ...                               // (critical code).
9 omp_unset_lock(&regionLock);      // UNLOCK.
10 ...
11 // Deallocate the lock.
12 omp_destroy_lock(&regionLock);
```

You *could* implement your own critical region this way, although it is easier to use `#pragma omp critical`.



# Programming locks

Note there is no **explicit** link between the lock `regionLock`, and the critical region of code, or data structure, that it is trying to protect.

It is down to the **programmer** to correctly link each lock with its associated block of critical code, or data structure.

This gives greater **flexibility**, but also greater scope for programming errors.

- Could use a `struct` or `class` to keep the lock with the data it is protecting, with the lock `private/protected`.

## Lock mistakes (1): Forgetting to lock()

```
1 lock_t regionLock;  
2 ...  
3 //regionLock.lock(); // Forgot to lock()!  
4 ...  
5 ... // Critical code  
6 ...  
7 regionLock.unlock();
```

This is precisely the situation we were trying to avoid!

- **All** threads enter the critical region.
- **Race conditions** become a possibility.

`unlock()` will have no effect, except possibly a small performance overhead<sup>1</sup>.

---

<sup>1</sup>Generally, this depends on the API: In C++11, attempting to unlock a `std::mutex` that is **not** locked leads to undefined behaviour.

## Lock mistakes (2): Forgetting to unlock()

```
1 lock_t regionLock;  
2 ...  
3 regionLock.lock();  
4 ...  
5 ... // Critical code  
6 ...  
7 //regionLock.unlock(); // Forgot to unlock!  
8
```

- The first thread **exclusively** enters the critical region.
- It never **releases** the lock.
- Therefore no other thread can **acquire** the lock.
- **All other threads remain idle at lock().**

# RAII = Resource Acquisition Is Initialisation.

This second mistake is easier to make than it seems:

- The critical code may throw an **exception** (C++/Java).
- A `break` or `continue` command may jump over `unlock()`.

May be support for locks that automatically release when they leave their scope.

- If defined at start of a routine, automatically released at end of routine **however it reached there**.
- e.g. `std::lock_guard<std::mutex>` in C++11.

This mechanism is generally known as RAII, for Resource Acquisition Is Initialisation.

# Multiple locks

Code on Minerva: `multipleLockCopy.c`

Suppose we want to copy randomly-selected elements of an array data of size  $N$  to another randomly-selected element.


- Decide to use a **lock** to control data writes.

```
1 #pragma omp parallel for
2 for( n=0; n<N; n++ )
3 {
4     i = rand() % N;
5     j = rand() % N;
6
7     omp_set_lock( &entireLock );    // Lock.
8     data[j] = data[i];              // Safe copy.
9     omp_unset_lock( &entireLock );  // Unlock.
10 }
```

# Multiple locks for memory access


This works, but is very inefficient.

- Only one thread can access the array data at a time.
- Even though only writing to one value.



Better to use **multiple locks** spanning the array:

- Different threads can write to different regions of the array **simultaneously**.
- Less **idle time** spent waiting for a lock to be released.



Using multiple locks is measurably faster (*try the code*):

```
1 omp_lock_t partialLocks[M];
2
3 // Initialise M locks near start of code.
4 ...
5 // Identify lock for writing to array.
6 int lock = M*j/N;
7 omp_set_lock( &partialLocks[lock] );
8 data[j] = data[i];
9 omp_unset_lock( &partialLocks[lock] );
10 ...
11 // Destroy all locks at end of code.
```

Note we only lock for the **write** to element  $j$

- Recall that just reading does **not** invoke a data race.

# Multiple locks for swapping

Code on Minerva: `multipleLockSwap.c`

Suppose now we want to **swap** elements *i* and *j*.

- Want to protect **each write** during the swap.

If access to the whole array was governed by a single lock, this would be straightforward to implement:

```
1  omp_set_lock( &entireLock );  
2  
3  // Writes to both data[i] and data[j].  
4  float temp = data[i];  
5  data[i] = data[j];  
6  data[j] = temp;  
7  
8  omp_unset_lock( &entireLock );
```

However, performance would again be poor.



## Multiple locks for swapping

We might think of using **two** locks, one for each region of the array being written to:

```
1  int lock_i = M*i/N;  
2  int lock_j = M*j/N;  
3  
4  omp_set_lock( &partialLocks[lock_i] );  
5  omp_set_lock( &partialLocks[lock_j] );  
6  
7  float temp = data[i];  
8  data[i] = data[j];  
9  data[j] = temp;  
10  
11 omp_unset_lock( &partialLocks[lock_i] );  
12 omp_unset_lock( &partialLocks[lock_j] );
```

Try this out!

## Why does this fail? – deadlock

Suppose one thread tries to lock `lock_i` then `lock_j`, and **simultaneously** another tries to lock `lock_j` *then* `lock_i`.

- Thread 1 **owns** `lock_i`, **waits** for `lock_j` to be released.
- Thread 2 **owns** `lock_j`, **waits** for `lock_i` to be released.

Since each thread is waiting for the other lock, they will never release the lock they own. **They will both wait forever.**

Threads waiting for synchronisation events that will never occur is known as **deadlock**.

The 'forgetting to `unlock()`' example earlier is also **deadlock**.

# Nested critical regions

Code on Minerva: `nestedCriticalRegion.c`

Another problem is when `lock_i==lock_j`. A simpler example where this occurs is for **nested critical regions**:

```
1 // Outer critical region
2 omp_set_lock( &lock )
3
4 // Inner critical region
5 omp_set_lock( &lock );
6 ...
7 omp_unset_lock( &lock ); // End of inner region
8
9 omp_unset_lock( &lock ); // End of outer region
```

In OpenMP, this will also **deadlock**.

- A thread that **owns** a lock cannot **re-acquire** the lock.

# Nested `#pragma omp critical`

OpenMP does **not** allow nested critical regions:

```
1 #pragma omp critical
2 {
3     ...
4     #pragma omp critical
5     {
6         ...
7     }
8 }
```

...will **not** compile.

- The **same** lock is being used by **both** critical sections.
- The same problem as in the previous slide.

## Named critical regions

This can be resolved by using **named** critical regions:

```
1 #pragma omp critical (OUTER)
2 {
3     ...
4     #pragma omp critical (INNER)
5     {
6         ...
7     }
8 }
```

- OUTER and INNER are user-defined labels.
- Each unique label corresponds to a unique lock.
- You are implicitly using a different lock for each critical region, so no thread tries to re-acquire a lock it already owns.

## Reacquiring locks

OpenMP code **deadlocks** if a thread tries to reacquire a lock it already owns.

- Just as it does not support nested critical regions.

This is primarily due to performance considerations:

- Would incur an execution overhead that would not always be necessary.

Not all parallel/concurrent APIs impose the same requirement - need to check the documentation!

- e.g. For C++11's `std::mutex`, the behaviour is undefined.
- Should also check documentation if attempting to `unlock` a lock that was **not** acquired.

# Summary of shared memory systems

Lecture	Content	Key points
2	Architectures and OpenMP	Cache coherency; false sharing; kernel <i>versus</i> user threads.
3	Data parallel problems	Fork-join construct; non-determinism; embarrassingly parallel problems.
4	Theory	Amdahl's law (strong scaling); Gustafson-Barsis law (weak scaling).
5	Data races	Loop parallelism; data dependencies.
6	Critical regions	Thread coordination; thread safety; serialisation; atomics.
7	Locks/mutexes	Performance costs for locks; dead-lock; named critical regions.

## Next lecture

Next time we will start to look at **distributed memory systems**:

- Multiple **processes** with their own **heap memory**.
- Examples: Clusters, super-computers.

Not surprisingly, data races are *not* an issue, but many of the other aspects we have covered *are*:

- Non-determinism, scaling, deadlock, data and loop parallelism, . . .