

# XJCO3221 Parallel Computation

Peter Jimack

University of Leeds

Lecture 16: GPU memory types

# Previous lectures

In the last lecture we saw how to program a simple GPU application:

- **Allocate** memory on the **device** (the GPU).
- **Copy host** (CPU) data to the device.
- Build and execute a **kernel** on the device that performs the computation.
  - Performed by many **work items** (*threads*).
  - Arranged into **work groups** of programmable size.
  - Can arrange work items in 1, 2 or 3 dimensions; the **NDRange** (= *n*-dimensional range).
- **Copy** the result back from the device to the host.

# This lecture

In the vector addition kernel, all of the arguments had the prefix `__global`:

```
1 __kernel
2 void kernel(__global float *a, __global float *b, ...)
3 {
4     // Kernel body.
5 }
```

In this lecture we will see what this means:

- The different **memory types** available to a GPU.
- How and when to use them.
- Performance issues related to **register overflow**.

# GPU memory

GPUs are designed as **high throughput devices**.

- **Many** threads (100's, 1000's, ...) execute simultaneously.
- By contrast, CPUs are **latency reducing** architectures, *i.e.* fast memory access by use of caches<sup>1</sup>, instruction level parallelism [*cf. Lectures 1, 2*], *etc.*

To maximise throughput, GPUs have multiple **memory types**:

- Architecture varies greatly between, *and within*, vendors.
- Performance would ideally be optimised for **each** GPU on which the code may be deployed.

---

<sup>1</sup>Although modern GPUs increasingly also have memory caches.

## Shared virtual, or ‘unified’, memory

In this module we treat CPU and GPU memory as **separate**.

- Typical of early GPU architectures.

Increasingly, CPU and GPU memory are presented to the programmer as **unified**<sup>1</sup>.

- API decides whether CPU or GPU memory is allocated.
- Integrated GPUs may share **physical** memory with the CPU.
- Supported from OpenCL 2.0, CUDA 4.0.

We will not consider unified memory in this module.

---

<sup>1</sup>See e.g. Wilt, *The CUDA Handbook* (Addison-Wesley, 2013); Han and Sharma, *Learn CUDA Programming* (Packt, 2019).

# Memory coalescing

When copying from e.g. global to local memory, **adjacent** threads will often access **adjacent** memory locations.

GPUs detect and optimise for this by **memory coalescing**:

- **Coalesces** multiple small memory accesses into a single large memory access — much faster.
- To exploit this, programmers can ensure nearby threads access nearby memory locations wherever possible.
- For 2D/3D data, can use a technique known as **tiling**<sup>1</sup>.

You are not expected to optimise your code for memory coalescence in this module.

---

<sup>1</sup>Rauber and Rünger, *Parallel Programming* (Springer, 2012). Tiling is also used to optimise cache access in CPUs.

# Memory types

GPUs typically contain 4 different types of memory.

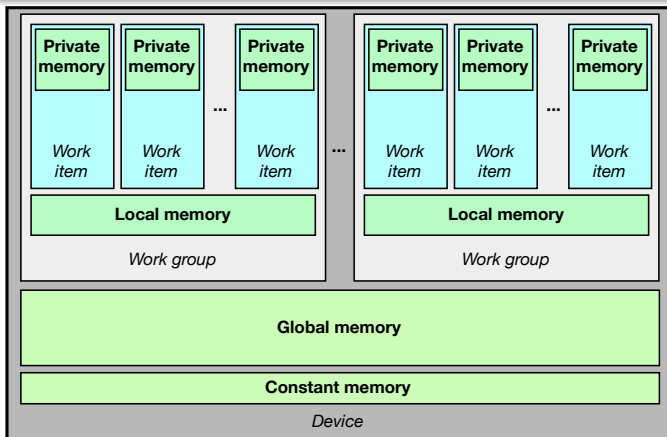
<b>Global</b>	Accessible to <b>all</b> work items in <b>all</b> work groups.
<b>Local</b> <sup>1</sup>	Shared by work items within <b>one work group only</b> .
<b>Private</b>	Accessible to a <b>single work item only</b> .
<b>Constant</b>	Global memory optimised for <b>read-only</b> operations. Not available on all GPUs.

These are **disjoint** - it is **not** allowed to cast one address space to another.

---

<sup>1</sup>In CUDA and Nvidia devices, *local* memory is referred to as *shared*.

# GPU memory types<sup>1</sup>



<sup>1</sup>After Kaeli *et al.*, *Heterogeneous computing with OpenCL 2.0* (Morgan-Kaufman, 2015).



## Analogy with CPU memory

Notice there is a **loose** analogy with CPU memory types:

CPU	GPU	Similarity
Shared memory	Global memory	Accessible by all processing units [ <b>threads</b> (CPU); <b>work items</b> (GPU)]
Distributed memory	Local memory	Only accessible to one <b>process</b> (CPU); <b>work group</b> (GPU).

To send data between work groups, must use global memory (or the host); a form of **communication**.

Cannot **directly** send data between the local memory of different work groups.

- The analogy with distributed memory CPUs breaks down.

# Allocating device memory

When allocating buffer memory on a device, it is placed in **global** memory.

For example, to allocate a global buffer called `device_array` capable of storing `N` floats:

```
1 cl_mem device_array= clCreateBuffer(context,  
    CL_MEM_READ_ONLY,N*sizeof(float),...);
```

Note that `CL_MEM_READ_ONLY` does **not** make it constant memory.

- Still allocated in the device's **global** memory.

## Read and write buffer flags

For OpenCL the buffer type should be specified as read-only, write-only, or read-write, by using one of the following flags:

<code>CL_MEM_READ_ONLY</code>	The buffer should only be read from.
<code>CL_MEM_WRITE_ONLY</code>	The buffer should only be written to.
<code>CL_MEM_READ_WRITE</code>	Both read and write allowed.

- Hint to allow **optimisations** by the runtime system.
- The default is `CL_MEM_READ_WRITE`.
- These refer to the **device** accessibility, *i.e.* **from inside kernels**, *not* from the host.

# Memory type 1. Global memory

## Global memory

Accessible to all processing units, but bandwidth is **slower** compared to the other memory types. Typically cached in modern GPUs, but not in older devices.

This is the memory type we have used already (*cf.* Lecture 15).

- Convenient from a programming perspective.
- Generally poor performance, although typically still faster than host-device communication.

# Using global memory in OpenCL<sup>1</sup>

Allocate global memory using `clCreateBuffer()`:

```
1 cl_mem device_buffer = clCreateBuffer(context, flags,
    size, ...);
```

In the kernel, prepend `__global` before the pointer(s):

```
1 __kernel
2 void vectorAdd( __global float *a, __global float *b,
    __global float *c )
3 {
4     int gid = get_global_id(0);
5     c[gid] = a[gid] + b[gid];
6 }
```

---

<sup>1</sup>In CUDA: `clCreateBuffer()` → `cudaMalloc()`; no `__global` specifier.

## Memory type 2. Local memory

### Local memory

Accessible to all work items in a work group, but **not** between groups. Much faster than global memory.

Typically used as a **scratch space** for calculations involving more than one work item in a group.

- Use for intermediate calculations.
- Place final answer in global memory.

In practice, this also requires **synchronisation** which is next lecture's topic, so will see an example of local memory then.

## Static local arrays

To use local memory in a kernel, simply place `__local` before the variable<sup>1</sup>:

```
1 __kernel
2 void kernel(__global int *device_array)
3 {
4     __local int temp[128];
5     // Calculations involving temp.
6     ...
7     // Place final answer in device_array.
8 }
```

However, this is **static allocation** - the size of the array must be known **when the kernel is built**.

---

<sup>1</sup>In CUDA: `__local` → `__shared__`.

## Dynamic local arrays

To create **dynamic local** arrays, declare it as a **kernel argument** with the specifier `__local`:

```
1 __kernel
2 void kernel(__local int *temp, __global int *device_a)
3 {
4     // Calculations involving temp.
5     ...
6     // Place final answer in device_a.
7 }
```

Then when setting kernel arguments, specify the size **but set the pointer to NULL**:

```
1 clSetKernelArg(kernel, 0, N*sizeof(int), NULL);
```



## Memory type 3. Private memory

### Private memory

Only accessible to each work item. *Very fast access.*

In practice, private memory is almost always implemented in hardware as **registers**:

- Small amount of memory that can be accessed quickly.
- Faster access than even local memory.
- **Automatic storage duration**, *i.e.* deallocated at the end of the kernel (or enclosing code block).

## Using private memory

There is no specifier for private memory (*i.e.* no `__private`).

- Variables declared within a kernel **default** to private.

For instance, in this kernel ...

```
1 __kernel
2 void kernel(__global float *array)
3 {
4     int gid = get_global_id(0);
5     ...
6 }
```

...the variable `gid` is treated as private memory.

## Private kernel arguments

Kernel arguments without a specifier are also treated as private.

In this example, `N` is treated as a private variable:

```
1 __kernel
2 void kernel(__global float *array, int N)
3 {
4     // Calculations involving N.
5     ...
6 }
```

The corresponding call to `setKernelArg()` would be:

```
1 int N=...;
2 clSetKernelArg(kernel, 1, sizeof(int), &N);
```

# Register overflow

Code on Minerva: `registerOverflow.c`, `registerOverflow.cl`, `helper.h`

Devices have a **fixed** amount of register memory. What happens if this is exceeded is device-dependent, but is usually one of:

- ① Private memory will ‘spill over’ into **global memory**.
- ② Fewer work groups will be launched simultaneously, resulting in an **under-utilisation** of available processing units.

Either mechanism reduces performance.

## Guidance

Kernels should be **small** (in the sense of low register usage) to limit the risk of register overflow.

## Memory type 4. Constant memory

### Constant memory

Accessible by all work items and work groups, but read only (by a kernel). Faster than global memory. Not available on all GPUs.

GPUs often have memory that is global in scope, but can only be read from within kernels:

- Known as **constant** memory.
- Can still be written to **by the host**.
- Originally to accelerate the mapping of **textures** to polygons ('texture' memory<sup>1</sup>).
- Typically much smaller than global memory, even if it exists.

---

<sup>1</sup>CUDA treats texture memory separately to other constant memory.

# Using constant memory<sup>1</sup>

For kernel arguments, use the `__constant` specifier:

```
1 __kernel
2 void kernel(__constant float *a,...)
3 { ... }
```

Initialise the array (*i.e.* copy from the host) and set the kernel argument as normal.

Variables within kernels can also be `__constant`, but must be specified at compile time:

```
1 __constant float pi = 3.1415926;
```

---

<sup>1</sup>In CUDA: Device data specified `__constant__`, and copy from host using `cudaMemcpyToSymbol()`.

## Summary and next lecture

Today we have look at the different **memory types** in GPUs:

- **Global** (slow), **local** (faster), **private** (very fast).
- Possibly also **constant** (faster than global).
- Private memory can **overflow**, resulting in performance loss.

The main use of **local** memory is to coordinate calculations between work items in a group. We will see an example next time when we look at **synchronisation** in GPUs.