

XJCO3221 Parallel Computation

Peter Jimack

University of Leeds

Lecture 17: Synchronisation

Previous lectures

Many of the previous lectures have mentioned **parallel synchronisation** in some form. However, there are many ways to synchronise:

- **Locks** in shared memory systems [*Lectures 6 and 7*].
- **Synchronisation barrier** at each level of a binary tree **reduction** [*Lecture 11*].
- **Blocking communication**, which affords a form of synchronisation, in distributed memory systems [*Lecture 9*].
- ...

Also recall that GPU's have multiple memory types, some of which can be viewed as *shared* (`__global`), and some which can be viewed as *distributed* (`__local`) [*Lecture 16*].

This lecture

In this lecture we will look at **synchronisation** on a GPU:

- How to synchronise **within** a work group.
- How to synchronise **between** work groups.

We will also see how the SIMD cores can potentially **reduce** or **improve** performance:

- Threads within a **subgroup** are **automatically synchronised**.
- Threads performing different calculations can lead to **divergence** and reduced performance.

Reminder: Scalar product

As an example, we will use the **scalar product** between two n -vectors **a** and **b**, as in Lecture 11.

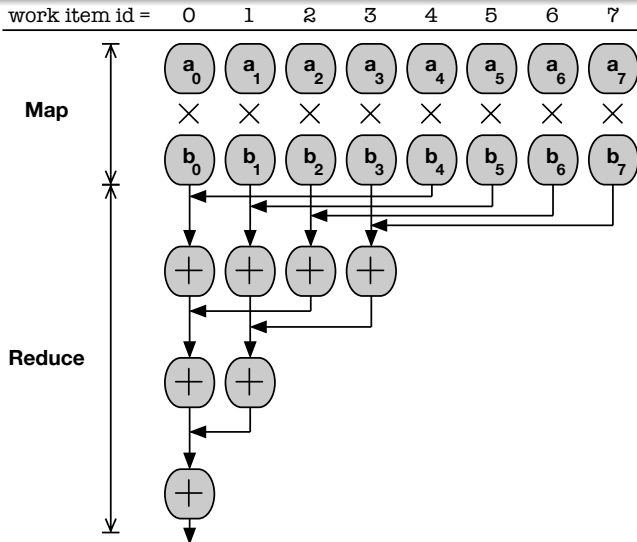
Written mathematically as (*indexing starting from 1*):

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=1}^n a_i b_i = a_1 b_1 + a_2 b_2 + \dots + a_n b_n$$

In serial CPU code (*indexing starting from 0*):

```
1 float dot = 0.0f;  
2 for( i=0; i<n; i++ ) dot += a[i] * b[i];
```

MapReduce pattern for $n = 8$



Reduction in local memory

First consider n equal to or less than the **work group size**.

Use **local memory** for the intermediate quantities [*Lecture 16*].

- **Faster** than global memory; allows **communication** within a work group.

Each work item copies $a[i]*b[i]$ to local memory first.

- Reduce using the binary tree pattern on the previous slide¹.
- Each addition performed by the work item with the *lower* i.d.
- Final result in work item with i.d. = 0 copied to the answer (in global memory).

¹Divide-by-two implemented by **compound bitwise right shift** operator '>>='.

Kernel code

Code on Minerva: `workGroupReduction.c`, `workGroupReduction.cl`, `helper.h`

```
1  __kernel
2  void reduceNoSync( __global float *device_a, __global
    float *device_b, __global float *dot, __local
    float *scratch )
3  {
4      int stride,
5          id      = get_local_id  (0),
6          groupSize = get_local_size(0);  // =work group
7
8      scratch[id] = device_a[id] * device_b[id];
9
10     for( stride=groupSize/2; stride>0; stride>>=1 )
11         if( id < stride )
12             scratch[id] += scratch[id+stride];
13
14     if(id==0) *dot = scratch[0];
15 }
```

Calling C-code

```
1 // float array of size 1 on device.
2 cl_mem device_dot = clCreateBuffer(...);
3
4 ... // Set kernel arguments 0, 1 and 2.
5 clSetKernelArg(kernel,3,N*sizeof(float),NULL);
6 // NULL => __local memory of given size.
7
8 // Add to the command queue.
9 size_t indexSpaceSize[1]={N}, workGroupSize[1]={N};
10 clEnqueueNDRangeKernel(queue,kernel,1,NULL,
    indexSpaceSize,workGroupSize,0,NULL,NULL);
11
12 // Get the result back to host float 'dot'.
13 float dot;
14 clEnqueueReadBuffer(queue,device_dot,CL_TRUE,0,sizeof(
    float),&dot,0,NULL,NULL);
```


Barriers

Without synchronisation, this reduction is *not* guaranteed to work on *all* systems.

Recall that **barriers** are points in code that no processing unit can leave until **all** units reach it [*c.f. Lecture 11*].

- `#pragma omp barrier` in OpenMP.
- `MPI_Barrier()` in MPI.

In OpenCL¹, a barrier **within a work group** is implemented as:

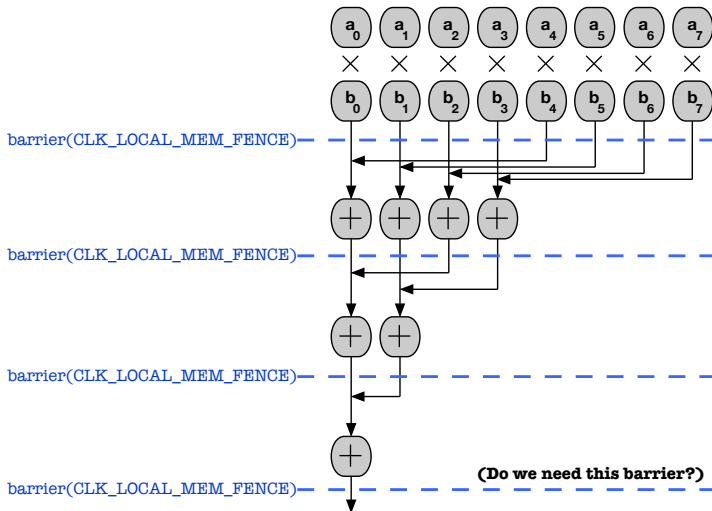
```
1 barrier(CLK_LOCAL_MEM_FENCE);
```

¹In CUDA: `__syncthreads()` synchronises within a **thread block**=work group.

Reduction with synchronisation

```
1 void reduceWithSync(...)    // Same arguments.
2 {
3     int id=..., groupSize=..., stride; // As before.
4
5     scratch[id] = device_a[id] * device_b[id];
6     barrier(CLK_LOCAL_MEM_FENCE);    // Sync.
7
8     for( stride=groupSize/2; stride>0; stride>>=1 )
9     {
10         if( id < stride )
11             scratch[id] += scratch[id+stride];
12
13         barrier(CLK_LOCAL_MEM_FENCE);    // Sync.
14     }
15
16     if(id==0) *dot = scratch[0];
17 }
```

Reduction with barrier(CLK_LOCAL_MEM_FENCE)



Problems larger than a single work group?

If we could synchronise **between** work groups, could use the same method as before:

- 1 Make device vectors and scratch **global**.
- 2 Replace **local** barriers with **global** barriers.

However, **no such global barrier exists**¹.

GPUs cannot synchronise between work groups/thread blocks²

¹`barrier(CLK_GLOBAL_MEM_FENCE)` *does* exist, but refers to *accesses* to global memory; it still only synchronises *within* a work group.

²Some modern GPUs support **cooperative groups** that allow synchronisation across multiple thread blocks; e.g. CUDA 9.0.


Warning!

You might see claims that it is possible to synchronise globally on any GPU by constantly **polling** a global memory location.

- *i.e.* work items constantly read/write to synchronise.

This may work, but **only for small problems**.

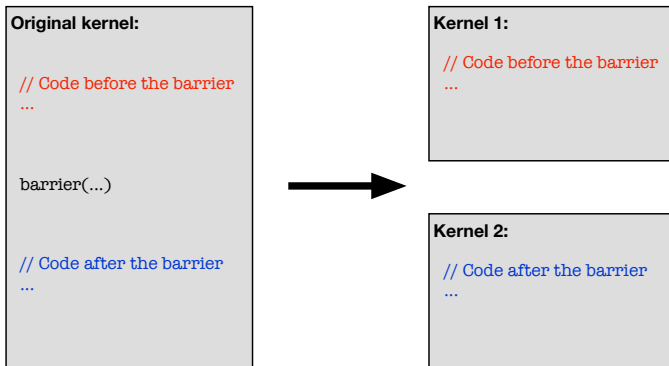
If there are too many work groups for the device, it **queues** them:



If they are not all on the device at the same time, it is **impossible to synchronise within one kernel** using this method.

Solution: Multiple kernels

The solution is to break the kernel at the barrier point into **multiple kernels** called consecutively:



This way kernel 1 completes before kernel 2 starts.

Reduction across work groups

It is possible to use this method for reduction¹:

- 1 Repeatedly call kernel that reduces an array of **partial sums** until less than maximum work group size.
- 2 Final kernel call to reduce these partial sums.

It is simpler (although less efficient) to use the CPU:

- 1 Each work group inserts its partial sum into a global array.
- 2 Final summation performed **on the host**.

This is conceptually similar to an MPI program performing final calculations on rank 0.

¹Wilt, *The CUDA handbook* (Addison-Wesley, 2013).

Subgroups (warp, wavefront, etc.)

Recall that GPUs are based on SIMD cores.

- Each core contains **multiple hardware threads** that perform the **same operation**.

In OpenCL, the number of **work items** (*i.e.* **threads**) simultaneously on a single SIMD core is known as a **subgroup**.

- **Smaller** than a work group.


The actual size is vendor specific. For example:

- Nvidia call them **warps**, each of which has 32 threads.
- AMD have 64-thread **wavefronts**.

Lockstep

The SIMD core applies the **same** operation to all items in the subgroup **simultaneously**. We say it advances in **lockstep**.

```
1 __kernel
2 void kernel(...)
3 {
4     int id = get_global_id(0);
5     float a, b, c;
6
7     a = 4*array[id];
8
9     b = a*a;
10
11    c = b + a;
12 }
```



Reduction with a subgroup

For reduction, this means that once the problem has been reduced to the size of a subgroup, there is **no longer any need for explicit synchronisation**¹:

```
1 __kernel
2 void reduce(...)
3 {
4     ...    // Start as before.
5
6     // Split the loop into two.
7     for(stride=group/2; stride>subgroup; stride>>=1) {
8         if(id<stride) scratch[id] += scratch[id+stride];
9         barrier(CLK_LOCAL_MEM_FENCE);    // Sync.
10    }
11    // See next slide ...
```

¹Wilt, *The CUDA handbook* (Addison-Wesley, 2013).

Final reduction

For the final reduction, exploit **lockstepping** by removing the **synchronisation**:

```
1 for(;stride>0;stride>>=1)
2 {
3     if(id<stride)
4         scratch[id] += scratch[id+stride];
5
6     // No barrier().
7 }
```

This avoids any overheads with calling `barrier()` (*i.e.* unnecessarily checking if all threads have reached this point, when we know they must have).

Divergence

Recall that SIMD cores can also be termed **SIMT** [*Lecture 14*]:

- Single Instruction stream, Multiple Threads.

This means it *is* possible to perform different operations with a subgroup, but they **remain in lockstep**.

- Only **one** distinct operation can be performed **at a time**.

This can lead to **serialisation** where operations are performed **one after the other**.

- Can lead to a severe performance penalty.


Code that leads to divergence

Suppose we want even-numbered work items to perform one operation, and odd-numbered items a different one¹:


```
1 __kernel
2 void kernel(...)
3 {
4     int id = get_global_id(0);
5
6     if( id%2 )
7         operation1;           // id odd.
8     else
9         operation2;           // id even.
10 }
```

¹Recall $i\%2==1$ for i odd, 0 for i even.

In this example the execution time is **double** what was expected:



For more operations the execution time increases further, e.g. a switch-case clause where every thread performs a different operation.



This is true **serialisation**.

Summary and next lecture

Today we have looked at **synchronisation**, focussing on GPUs.

- **Barriers** can synchronise within a work group.
- **Cannot** synchronise between work groups within a kernel - must split into separate kernels.
- Work items of threads within a **subgroups** execute in lockstep.
 - No need for explicit synchronisation mechanisms.
 - Can lead to **divergence** and reduced performance.

Next time we will look at **atomic** instructions, continuing what we started in Lecture 6.